

CS 489/698 Assignment 2

TA: Liyi Zhang (l392zhang@uwaterloo.ca)

Office hours: June 22 (Thursday), 1:00pm - 2:00pm, DC 3333

June 28 (Wednesday), 10:00am - 11:00am, DC 3333

Due date: June 30, 2023

Part 1: Bugs That Are Hard to Catch

Program analysis is used both by attackers and defenders to hunt for potential issues in the software. However, there are always code patterns that pose unique challenges for these tools to analyze. The goal of this part of the assignment is to help you understand the *limitations* of state-of-the-art program analysis tools and get a feeling on their advantages and disadvantages with some hands-on experience.

In short, your goal is to craft programs **with a single bug intentionally embedded** and check whether different program analysis tools can find the bug (evident by a program crash). If the program analysis tool fails to find the bug *within a computation bound*, you have found a potential limitation in this program analysis tool and will be awarded full marks for that component.

In this assignment, the computation bound is **15 minutes per analyzer \times 3 analyzers**:

- AFL++: as a representative fuzzer
- KLEE: as a representative symbolic executor
- SymCC: as a representative concolic executor (and hybrid fuzzer)

We will provide details on how to run these tools in later part of the assignment.

To confine the scope of the analysis and standardize the auto-grading process, we will NOT accept arbitrary programs for this assignment. Instead, you are encouraged to produce *minimal* programs and prepare a package for each program. The package should contain not only the program code but also information that bootstraps analysis and shows evidence that a bug indeed exists.

Specifically, each package MUST be prepared according to the **requirements and restrictions** below. Failing to do so will result in an invalid package that can't be used to score this assignment component.

- Each submitted package MUST follow the directory structure below:

```
<package>/
|-- main.c      // mandatory: the only code file to submit
|-- input/      // mandatory: the test suite with 100% gcov coverage
|   |-- <*>     // test case name can be arbitrary valid filename
|-- crash/      // mandatory: sample inputs that crash the program
|   |-- <*>     // sample name can be arbitrary valid filename
|-- README.md   // optional: an explanatory note on the code and embedded bug
```

- Include all source code of the program in **one and only one** `main.c` file. This `main.c` is the only code file you need to include in the package. DO NOT include `interface.h`.
- Only valid **C** programs are allowed. DO NOT code in C++.

- The size of `main.c` should NOT exceed 256KB (i.e., 256×2^{10} bytes).
- Your program can only invoke three library calls, all provided in `interface.h` available [here](#).
 - `ssize_t in(void *buffer, size_t count)`
which reads in at max `count` bytes from `stdin` and store then in `buffer`. The return value indicates the actual number of bytes read in or a negative number indicating failure.
 - `int out(const char *buffer)`
which prints the `buffer` string to `stdout`. The return value indicates the actual number of bytes written out.
 - `void abort(void)`
which forces a crash of program. Note that this is NOT the only way to crash a program.
- Your program can only take input from `stdin` using the provided `in()` function in `interface.h`. It should NOT take input from command line arguments nor environment variables. The size of input acquired from `stdin` should NOT exceed 1024 bytes.
- Your program MUST be *compatible* with **all** program analysis tools used in this assignment without special modification to these tools. In other words, your program can be analyzed using the tool invocation command provided in later part of the assignment.
- Your program should have **one and only one bug** that is intentionally planted. If any of the tool finds a crash in your program, even the crash is not caused by the intended bug, the tool is considered successful and this package cannot be used to claim victory over that tool.
- Provide a set of test cases that achieves 100% coverage (see details on `gcov` below).
 - Each test case should NOT crash the program, i.e., exiting with a non-zero status code.
 - Each test case should complete its execution in 10 seconds.
 - Each test case should NOT exceed 1024 bytes in size.
- Provide **at least one** sample input that can cause the program to crash by triggering the planted bug — this is to provide evidence on the existence of the bug.
- Avoid using features that are known limitations for most program analysis tools, including:
 - DO NOT use any other library functions (including `libc` functions). The only permitted library routines are given in the header file `interface.h`.
 - DO NOT use floating-point operations in your program.
 - DO NOT use inline assemblies in your program.
 - DO NOT use multi-threading. The entire program logic must be able to execute end-to-end in a single process and a single thread.

Each package will be analyzed by all program analysis tools. For any tool *X*, if the tool finds a bug (even not the planted one), it is considered a success. Otherwise, the bug has managed to “evade” the detection from tool *X* and the package will be counted towards the scoring of tool *X* evasion.

A maximum of 10 packages can be submitted. You will have full points for tool *X* evasion as long as you have one package that “evaded” the detection from tool *X*. Below is the more verbose criteria:

- [10 pts] at least one package that evades AFL++
- [10 pts] at least one package that evades KLEE
- [10 pts] at least one package that evades SymCC
- [10 pts] one package that evades them all

If you are confident that one package can be used to score all four components, feel free to submit one package only. However, to be on the safe side, it is highly advised to submit multiple packages with different strategies to confuse these program analysis tools.

1.0 Environment preparation

We provide you with a Virtual Machine image to help you with the tool setup. This VM image is provisioned by the [Vagrant](#) tool using the [Vagrantfile](#) available in [this GitHub repository](#). If you are interested, feel free to clone the repository and “[vagrant up](#)” from there. Otherwise, you can setup the environment by importing the VirtualBox image follow the instructions below:

1. Install the free [VirtualBox \(Version 7.0.8\)](#) software.
2. Download the provided [virtual machine file](#) for this part.
3. Double-click on the downloaded virtual machine `.ova` file, which will open in VirtualBox. Complete importing appliance without changing the default settings.
4. Start the provided virtual machine in VirtualBox
After the virtual machine boots, you will see a login screen.
Both username and password are `vagrant`.
5. User `vagrant` has `sudo` rights, without password.
6. We aim to provide all necessary scripts and samples under directory `/home/vagrant/workdir`. However, you are free to install anything you need to complete this assignment. Just be aware that when your submitted packages are evaluated, they are evaluated on a freshly provisioned virtual machine image.
7. The VM image is a command line interface by default.
You do not have a desktop environment (although you may choose to install one via `apt`).
8. Network Configuration (for troubleshooting):
 - Open the VirtualBox software. In the menu bar, choose "File" -> "Tools" -> "Network Manager".
 - Select the "NAT Networks" tab.
 - Click on the "Create" button to create a new network named "NatNetwork" with "IPv4 Prefix" set to 10.0.2.0/24.

NOTE: This VM is also the evaluation environment for grading your packages. The VM will be run on a server machine with Intel Xeon Silver 4214R CPU @ 2.40GHz

1.1 Coverage tracking with gcov

[gcov](#) is a tool you can use in conjunction with `gcc` to test code coverage in your programs. In a nutshell, it tracks the the portion of code that is “covered” by a concrete execution at runtime and aggregates the coverage results from multiple runs to produce a final coverage report.

You can use the [run-gcov.sh](#) script to check the coverage of your package. The script is available to you, as a reference implementation, under `/home/vagrant/workdir/run-gcov.sh` in the provided VM. In general, the script performs the following steps:

```
// Step 1: compile your code with gcov instrumentations
$ gcc -fprofile-arcs -ftest-coverage -g main.c -o main

// Step 2: execute each of your input test case
$ for test in input/*; do main < ${test}; done

// Step 3: collect and print coverage
$ gcov -o ./ -n main.c
```

If you read a 100% coverage in your console, it means your test suite (provided under the `input/` directory) has achieved complete coverage in `gcov`’s perspective.

1.2 Fuzzing with AFL++

We use the open-source [AFL++ fuzzer](#), in particular, version 4.07c which is the most up-to-date stable release of AFL++. You may want to read a bit of details on the [project page](#) about AFL++ and fuzzing in general.

AFL++ is provisioned into the VM as a Docker image tagged as `afl`. Once in the VM, you can use the following command to run the Docker image, get an interactive shell, and explore around.

```
docker run \
  --tty --interactive \
  --volume <path-to-your-package>:/test \
  --workdir /test \
  --rm afl \
  bash
```

You can use the [run-afl.sh](#) script to fuzz your package. The script is available to you, as a reference implementation, under `/home/vagrant/workdir/run-afl.sh` in the provided VM. In general, the script runs-off the `afl` Docker image and performs the following steps:

```
// Step 1: compile your code with afl instrumentations
$ afl-cc main.c -o main

// Step 2: start fuzzing your code
$ afl-fuzz -i input -o output -- main
```

The output of the AFL++ fuzzing results are stored in the `output` directory.

1.3 Symbolic reasoning with KLEE

We use the open-source [KLEE symbolic executor](#), in particular, version v3.0 which is the most up-to-date stable release of KLEE. You may want to read a bit of details on the [project page](#) about KLEE and symbolic execution in general.

KLEE is provisioned into the VM as a Docker image tagged as `klee`. Once in the VM, you can use the following command to run the Docker image, get an interactive shell, and explore around.

```
docker run \
  --tty --interactive \
  --volume <path-to-your-package>:/test \
  --workdir /test \
  --rm klee \
  bash
```

You can use the [run-klee.sh](#) script to symbolically execute your package. The script is available to you, as a reference implementation, under `/home/vagrant/workdir/run-klee.sh` in the provided VM. In general, the script runs-off the `klee` Docker image and performs the following steps:

```
// Step 1: compile your code into LLVM IR
$ clang -emit-llvm -g -O0 -c main.c -o main.bc

// Step 2: symbolically explore all paths in the program
// NOTE: it caps the symbolic input from stdin to at max 1024 bytes
$ klee --libc=klee --posix-runtime --output-dir=output main.bc -sym-stdin 1024
```

The output of the KLEE execution results are stored in the `output` directory. The test cases generated by KLEE can be parsed by the [ktest-tool](#).

1.4 Concolic execution with SymCC

We use the open-source [SymCC concolic executor](#), in particular, commit 77bb971 which is the most up-to-date stable commit of SymCC. You may want to read a bit of details on the [project page](#) about SymCC, concolic execution, and hybrid fuzzing in general.

SymCC is provisioned into the VM as a Docker image tagged as `symcc`. Once in the VM, you can use the following command to run the Docker image, get an interactive shell, and explore around.

```
docker run \  
  --tty --interactive \  
  --volume <path-to-your-package>:/test \  
  --workdir /test \  
  --rm symcc \  
  bash
```

You can use the [run-symcc.sh](#) script to fuzz your package with a concolic executor. The script is available to you, as a reference implementation, under `/home/vagrant/workdir/run-symcc.sh` in the provided VM. In general, the script runs-off the `symcc` Docker image and performs the following steps:

```
// Step 1: compile your code with afl instrumentations  
$ /afl/afl-clang main.c -o main-afl  
  
// Step 2: compile your code with symcc instrumentations  
$ symcc main.c main.c -o main-sym  
  
// Step 3: start fuzzing your code with naive afl in the background  
$ screen -dmS afl -- \  
  /afl/afl-fuzz -M afl-0 -i input -o output -- main-afl  
  
// Step 4: wait a while for the fuzzer to initialize the directories  
$ sleep 5  
  
// Step 5: invoke the concolic executor to find more seeds  
$ symcc_fuzzing_helper -v -o output -a afl-0 -n symcc -- main-sym
```

The output of the SymCC hybrid fuzzing results are stored in the `output` directory.

1.5 Example

The provided VM image contains a sample package under directory `/home/vagrant/workdir/pkg-sample` to illustrate how a package should look like, with the addition of `interface.h` and `.gitignore` which shouldn't be submitted. The code can also be found on the [repository](#).

You can test out the sample package inside the VM via:

```
$ cd /home/vagrant/workdir  
$ ./run-gcov.sh pkg-sample  
$ ./run-afl.sh pkg-sample  
$ ./run-klée.sh pkg-sample  
$ ./run-symcc.sh pkg-sample
```

The output should obviously show that 1) this package does not provide 100% code coverage and 2) it has a bug that is found by all three tools.

Despite that this package sample cannot “evade” any program analysis tool, feel free to duplicate this template package to bootstrap your package preparation that can eventually “evade” the tools.

Part 2: Escaping seccomp-based Sandboxes

2.1 Environment preparation

Before starting this part, you need to set up the test environment. Please follow the instructions below:

1. Install the free [VirtualBox \(Version 7.0.8\)](#) software.
2. Download the provided [virtual machine file](#) for this part.
3. Network Configuration:
 - Open the VirtualBox software. In the menu bar, choose "File" -> "Tools" -> "Network Manager".
 - Select the "NAT Networks" tab.
 - Click on the "Create" button to create a new network named "NatNetwork" with "IPv4 Prefix" set to 10.0.2.0/24.
4. Double-click on the downloaded virtual machine .ova file, which will open in VirtualBox. Complete importing appliance without changing the default settings.
5. Start the provided virtual machine in VirtualBox.
6. After the virtual machine boots, you will see a login screen for the user `seed`, enter `dees` as the password.
7. On the desktop, you will find a folder named `part-2` contains the required files for this part.

2.2 Tasks

In this part, we have provided three binary executable files: `sandbox1`, `sandbox2`, and `sandbox3`. Each of these programs implements a `seccomp`-based sandbox and employs distinct filtering rules. The programs are designed to execute external programs `exploit1`, `exploit2`, and `exploit3` respectively. Additionally, we have provided a `/home/seed/flag` file that stores a flag you required to get.

Your task is to write three exploit programs in `C`, namely `exploit1.c`, `exploit2.c`, and `exploit3.c`, and compile them into three binary executable files `exploit1`, `exploit2`, and `exploit3`, respectively, in such a way that when running the sandbox programs, the contents of the flag file are printed. In this part, you are expected to submit the three source code files of your exploit programs with a `Makefile`, such that we can get your compiled binary executable files by simply typing `"make"` in the terminal. Below is an example demonstrating how to test your code and run the sandbox programs:

```
// compile your exploit programs
// (should produces exploit1, exploit2, and exploit3)
$ make

// run the sandbox programs
// (your exploit programs should be in the same folder as the sandbox programs)
$ ./sandbox1
the flag should be printed here

$ ./sandbox2
the flag should be printed here

$ ./sandbox3
the flag should be printed here
```

NOTE: At the time of grading your submissions, we will replace the contents of the flag file. Therefore, you should not hardcode the contents of the flag file in your program.

Each `exploit` carries an equal weight of **10 pts**, making a total of 30 pts for Part 2.

Deliverables

You are required to hand in a single compressed `.zip` file, which should unzip into two directories, `part-1` and `part-2`:

- Directory `part-1/` contains up to 10 packages where each package is a directory itself. We will count the first 10 packages only alphabetically ordered if more than 10 packages are found.
- Directory `part-2/` contains four deliverables: **`exploit1.c`**, **`exploit2.c`**, **`exploit3.c`**, and **`Makefile`**, as well as one optional write-up (see below).

Submit your files using Assignment 2 drop box in LEARN.

Write-up. We use an auto-grader to check the code submitted for this assignment for both parts. While efficient, the auto-grader can only provide a binary pass/fail result, which rules out the possibility of awarding partial marks for each task. As a result, we also solicit a write-up submissions.

- For Part 1, the write-ups are optionally included in each package as `README.md` files.
- For Part 2, the write-up can include information about any of the exploits you attempted as long as you believe the information is relevant for the grading. The write-up can be in any format.

Typical things to be put in the write-up include:

- How the code you submitted is expected to work
- How you manage to get certain hardcoded information in the code
- Explanation on critical steps / algorithms in the code
- Any special situations the TAs need to be aware when running the code
- If you do not complete the full task, how far you have explored

On the other hand, if you are confident that all your code will work out of the box and can tolerate a zero score for any tasks on which the auto-grader fails to execute your code, you do not need to submit a write-up (or you can omit certain tasks in the write-up).